

# SET A

## 1. Kubernetes Service Types

**Question:** You want to expose an application running inside your Kubernetes cluster to external users using a stable external IP. Which Kubernetes Service type should you use?

- A. **ClusterIP**
- B. **NodePort**
- C. **LoadBalancer**
- D. **ExternalName**

**Answer:** C

**Explanation:**

- **LoadBalancer (C)** provisions an external load balancer (if supported by the cloud provider) and assigns a public IP for accessing the service.
  - **ClusterIP (A)** is the default and allows only internal access.
  - **NodePort (B)** exposes the service on a static port on each node, but without an external load balancer.
  - **ExternalName (D)** maps a service to an external DNS name, but does not expose the service via an IP.
- 

## 2. Kubernetes Pod Scheduling Constraints

**Question:** You need to ensure that a specific pod always runs on a node with a GPU for machine learning workloads. Which Kubernetes feature should you use?

- A. **Taints and Tolerations**
- B. **NodeSelector**
- C. **Pod Disruption Budget (PDB)**
- D. **Pod Priority Classes**

**Answer:** B

**Explanation:**

- **NodeSelector (B)** allows you to specify a node with specific labels (e.g., `gpu=true`) to schedule pods on nodes with GPUs.

- **Taints and Tolerations (A)** prevent workloads from being scheduled on specific nodes unless explicitly tolerated.
  - **PDB (C)** is used to define how many replicas of a pod can be disrupted during maintenance, not for scheduling.
  - **Pod Priority Classes (D)** influence scheduling but do not restrict a pod to specific nodes.
- 

### 3. Kubernetes Networking Model

**Question:** What is true about the Kubernetes networking model?

- A. **Each pod gets a unique IP address and can communicate with other pods without NAT.**
- B. **Pods communicate using host machine IPs, and port mapping is required.**
- C. **Containers within the same pod cannot talk to each other directly.**
- D. **All network policies are enabled by default for inter-pod communication.**

**Answer:** A

**Explanation:**

- **Kubernetes assigns a unique IP to each pod (A)**, and all pods in a cluster can communicate without NAT.
  - **Host machine IPs (B)** are used only in certain cases like NodePort services, not for pod-to-pod communication.
  - **Containers inside the same pod (C)** share the same network namespace and can communicate via `localhost`.
  - **Network policies (D)** are not enforced by default; they must be explicitly created.
- 

### 4. Kubernetes Rolling Update vs. Recreate

**Question:** What happens during a **rolling update** of a Deployment in Kubernetes?

- A. **All existing pods are terminated immediately and new ones are created.**
- B. **Pods are replaced one at a time to ensure no downtime.**
- C. **Only some replicas are restarted while others remain running.**
- D. **A new Deployment is created, and the old Deployment is deleted.**

**Answer:** B

**Explanation:**

- **Rolling updates (B)** ensure availability by replacing pods incrementally.
  - **Recreate strategy (A)** would terminate all pods first, leading to downtime.
  - **Option C** is partially correct but doesn't explicitly state a controlled, sequential rollout.
  - **Option D** is incorrect—Deployments are updated, not replaced.
- 

## 5. Kubernetes ConfigMap vs Secret

**Question:** How do ConfigMaps differ from Secrets in Kubernetes?

- A. **ConfigMaps are stored as base64-encoded data, whereas Secrets store plain text.**
- B. **Secrets are stored securely and can be encrypted, while ConfigMaps store non-sensitive configuration.**
- C. **Both ConfigMaps and Secrets require pods to restart for changes to take effect.**
- D. **ConfigMaps can only be mounted as environment variables, while Secrets cannot.**

**Answer:** B

**Explanation:**

- **Secrets are encoded and can be encrypted for security (B),** whereas ConfigMaps hold non-sensitive configuration.
  - **ConfigMaps store plain text (A is incorrect),** while Secrets are base64-encoded by default.
  - **Option C** is incorrect—pods don't always need a restart; some applications can reload configurations dynamically.
  - **Option D** is incorrect—both can be used as environment variables or mounted as files.
- 

## 6. Kubernetes Persistent Volume Access Modes

**Question:** Which Kubernetes **Persistent Volume (PV) access mode** allows multiple nodes to **read and write** to the same volume at the same time?

- A. **ReadWriteOnce (RWO)**
- B. **ReadOnlyMany (ROX)**
- C. **ReadWriteMany (RWX)**
- D. **ReadWriteOnly (RWO)**

**Answer:** C

**Explanation:**

- **ReadWriteMany (RWX) (C)** allows multiple nodes to read and write simultaneously.
  - **ReadWriteOnce (RWO) (A)** means only one node can mount the volume in read/write mode.
  - **ReadOnlyMany (ROX) (B)** allows multiple nodes to mount but only for read access.
  - **D is an incorrect option (no such mode as ReadWriteOnly).**
- 

## 7. Kubernetes Horizontal Pod Autoscaler (HPA)

**Question:** You want to **automatically scale** the number of pods in a deployment based on CPU utilization. Which Kubernetes feature should you use?

- A. **Cluster Autoscaler**
- B. **Vertical Pod Autoscaler (VPA)**
- C. **Horizontal Pod Autoscaler (HPA)**
- D. **PodDisruptionBudget (PDB)**

**Answer:** C

**Explanation:**

- **HPA (C)** automatically scales the number of pods based on resource utilization.
  - **Cluster Autoscaler (A)** scales nodes, not pods.
  - **Vertical Pod Autoscaler (B)** adjusts pod resource requests but does not scale replicas.
  - **PodDisruptionBudget (D)** controls disruption limits but does not handle autoscaling.
- 

## 8. Kubernetes PodSecurityPolicy

**Question:** You need to **prevent pods from running as root** in your Kubernetes cluster. Which feature should you use?

- A. **PodDisruptionBudget (PDB)**
- B. **SecurityContext in pod definition**
- C. **RBAC (Role-Based Access Control)**
- D. **ServiceAccount**

**Answer:** B

**Explanation:**

- **SecurityContext (B)** can enforce that pods do not run as root.
- **PodDisruptionBudget (A)** controls pod eviction, not security settings.

- **RBAC (C)** manages permissions but does not enforce security settings on pod execution.
  - **ServiceAccount (D)** assigns identities to pods but does not enforce runtime security.
- 

## 9. Kubernetes Init Containers

**Question:** What is the primary purpose of **Init Containers** in Kubernetes?

- A. To provide long-running background services within a pod.
- B. To execute setup tasks before the main containers start.
- C. To replace the main application container in case of failure.
- D. To handle persistent storage for the pod.

**Answer:** B

**Explanation:**

- **Init Containers (B)** run before the main application containers to perform initialization tasks like setting up configuration files.
  - **They do not run as background services (A)**—they terminate after completion.
  - **They do not replace failed containers (C).**
  - **They do not directly handle storage (D).**
- 

## 10. Kubernetes DaemonSet vs Deployment

**Question:** You need to ensure that a specific pod is deployed to **every node in the cluster**. Which Kubernetes resource should you use?

- A. **Deployment**
- B. **StatefulSet**
- C. **DaemonSet**
- D. **ReplicaSet**

**Answer:** C

**Explanation:**

- **DaemonSets (C)** ensure that a pod runs on every node in the cluster (e.g., log collectors, monitoring agents).
- **Deployments (A)** manage replicas but do not guarantee placement on every node.
- **StatefulSets (B)** are used for stateful applications with stable network identities.

- **ReplicaSets (D)** ensure a specific number of pod replicas but do not control node placement.
- 

## SET B

### 1. Kubernetes Network Policy

**Question:** You want to ensure that only a specific pod in **namespace A** can access a pod running in **namespace B**. Which Kubernetes feature should you use?

- A. **Ingress Controller**
- B. **RBAC (Role-Based Access Control)**
- C. **NetworkPolicy**
- D. **ServiceAccount**

**Answer:** C

**Explanation:**

- **NetworkPolicy (C)** controls traffic flow between pods at the network level and can restrict access between namespaces.
  - **Ingress Controller (A)** handles external traffic, not inter-pod communication.
  - **RBAC (B)** manages API access, not network security.
  - **ServiceAccount (D)** is used for API authentication, not network policies.
- 

### 2. Kubernetes StatefulSets

**Question:** What is the key difference between a **StatefulSet** and a **Deployment** in Kubernetes?

- A. **StatefulSets use a unique identity and persistent storage for each pod.**
- B. **Deployments always ensure that pods run on different nodes.**
- C. **StatefulSets only support one replica per pod definition.**
- D. **Deployments cannot automatically scale replicas.**

**Answer:** A

**Explanation:**

- **StatefulSets (A)** maintain a stable identity and persistent storage for each pod, making them ideal for databases and applications that require ordered deployment.
  - **Deployments (B)** do not guarantee pods will be placed on different nodes.
  - **StatefulSets allow multiple replicas (C is incorrect).**
  - **Deployments (D) support autoscaling with Horizontal Pod Autoscaler.**
- 

### 3. Kubernetes Readiness vs. Liveness Probes

**Question:** What is the main purpose of a **readiness probe** in a Kubernetes pod?

- A. To check if a pod is running and ready to receive traffic.
- B. To restart the pod if it becomes unresponsive.
- C. To update DNS entries for the pod.
- D. To determine if a pod has been scheduled on a node.

**Answer:** A

**Explanation:**

- **Readiness probes (A)** indicate if a pod is ready to serve traffic.
  - **Liveness probes (B)** determine if a pod is healthy or needs a restart.
  - **DNS updates (C) happen automatically in Kubernetes.**
  - **Pod scheduling (D) is determined by the scheduler, not probes.**
- 

### 4. Kubernetes API Resource Hierarchy

**Question:** Which Kubernetes resource is the **parent** of a ReplicaSet?

- A. Pod
- B. Deployment
- C. DaemonSet
- D. ConfigMap

**Answer:** B

**Explanation:**

- **Deployments (B)** create and manage **ReplicaSets**, which in turn manage **pods**.
- **DaemonSets (C)** ensure that a pod runs on every node.
- **ConfigMaps (D)** store configuration but are not part of pod scheduling.

---

## 5. Kubernetes Job vs. CronJob

**Question:** You need to run a **one-time** data migration task in Kubernetes. Which resource should you use?

- A. **Job**
- B. **CronJob**
- C. **StatefulSet**
- D. **Deployment**

**Answer:** A

**Explanation:**

- **Jobs (A)** are for one-time batch tasks.
  - **CronJobs (B)** schedule recurring jobs.
  - **StatefulSets (C)** manage persistent applications.
  - **Deployments (D)** are for long-running applications.
- 

## 6. Kubernetes Service Discovery

**Question:** How do Kubernetes pods discover and communicate with services inside a cluster?

- A. **Using built-in DNS resolution for Service names.**
- B. **Manually assigning static IP addresses to pods.**
- C. **By using host-based routing configured in ConfigMaps.**
- D. **By exposing each pod with a NodePort.**

**Answer:** A

**Explanation:**

- Kubernetes **automatically assigns DNS names** to services, allowing pods to resolve service names without needing static IPs.
  - **Static IPs (B)** are not used since pods are ephemeral.
  - **ConfigMaps (C)** store configurations but do not control networking.
  - **NodePort (D)** exposes services externally but is not required for internal pod communication.
-

## 7. Kubernetes RBAC Roles

**Question:** Which **RBAC resource** is used to restrict access to **only specific resources** in a namespace?

- A. **ClusterRole**
- B. **Role**
- C. **RoleBinding**
- D. **PodSecurityPolicy**

**Answer:** B

**Explanation:**

- **Roles (B)** apply permissions within a specific namespace.
  - **ClusterRoles (A)** grant permissions cluster-wide.
  - **RoleBindings (C)** assign roles to users but do not define permissions.
  - **PodSecurityPolicy (D)** is used for security settings, not access control.
- 

## 8. Kubernetes Horizontal Pod Autoscaler (HPA) Metrics

**Question:** What metric does the **Horizontal Pod Autoscaler (HPA)** use by default to scale pods?

- A. **Memory utilization**
- B. **CPU utilization**
- C. **Disk I/O**
- D. **Pod restarts**

**Answer:** B

**Explanation:**

- **HPA scales based on CPU utilization (B) by default.**
  - **Memory utilization (A)** requires custom metrics.
  - **Disk I/O (C) and pod restarts (D)** are not used for autoscaling.
-

## 9. Kubernetes Cluster Autoscaler

**Question:** You need to automatically **add and remove nodes** from a Kubernetes cluster based on pod scheduling needs. Which feature should you use?

- A. **Cluster Autoscaler**
- B. **Horizontal Pod Autoscaler**
- C. **Vertical Pod Autoscaler**
- D. **NodeSelector**

**Answer:** A

**Explanation:**

- **Cluster Autoscaler (A)** adds/removes worker nodes based on demand.
  - **HPA (B)** scales pods, not nodes.
  - **VPA (C)** adjusts pod resource requests.
  - **NodeSelector (D)** assigns pods to specific nodes but does not scale them.
- 

## 10. Kubernetes Ingress Controllers

**Question:** You want to route incoming HTTP requests based on path rules (e.g., `/api` to one service, `/web` to another). Which resource should you use?

- A. **Service of type LoadBalancer**
- B. **Ingress**
- C. **NetworkPolicy**
- D. **ConfigMap**

**Answer:** B

**Explanation:**

- **Ingress (B)** handles HTTP routing based on host/path rules.
  - **LoadBalancer (A)** exposes a single backend service.
  - **NetworkPolicies (C)** control internal networking but do not handle HTTP routing.
  - **ConfigMaps (D)** store configurations but do not handle traffic routing.
-

## SET C

### 1. Kubernetes Pod-to-Pod Communication

**Question:** By default, how do pods communicate with each other in Kubernetes?

- A. Pods in the same namespace can communicate directly, but pods in different namespaces require NetworkPolicies.
- B. All pods can communicate with each other across nodes and namespaces without any restrictions.
- C. Pods communicate using NAT (Network Address Translation).
- D. Pods must explicitly expose a service to communicate with each other.

**Answer:** B

**Explanation:**

- Kubernetes uses a flat network model where all pods can communicate without NAT (B).
  - NetworkPolicies (A) must be explicitly defined to restrict traffic.
  - Pods communicate without NAT (C) because they are assigned unique IPs.
  - Services (D) provide stable access but are not required for pod-to-pod communication.
- 

### 2. Kubernetes Admission Controllers

**Question:** Which admission controller can prevent pods from being created if they do not have specific labels?

- A. ValidatingWebhook
- B. MutatingWebhook
- C. ResourceQuota
- D. PodSecurityAdmission

**Answer:** A

**Explanation:**

- ValidatingWebhook (A) can block resource creation based on policies (e.g., missing labels).
- MutatingWebhook (B) modifies requests but does not block them.
- ResourceQuota (C) enforces limits on resources but does not validate labels.
- PodSecurityAdmission (D) enforces pod security standards but does not check labels.

---

### 3. Kubernetes ServiceAccount Tokens

**Question:** What is the primary purpose of a **ServiceAccount** in Kubernetes?

- A. To provide network identity to a pod for inter-pod communication.
- B. To grant API access to pods running inside the cluster.
- C. To create persistent storage volumes for stateful workloads.
- D. To assign a unique hostname to each pod.

**Answer:** B

**Explanation:**

- **ServiceAccounts (B)** allow pods to authenticate to the Kubernetes API.
  - **Pods communicate via networking, not identity (A).**
  - **Storage (C)** is managed via PersistentVolumes.
  - **Hostnames (D)** are assigned based on pod DNS configurations.
- 

### 4. Kubernetes Persistent Volumes

**Question:** Which **StorageClass** reclaim policy allows Kubernetes to automatically delete the PersistentVolume when its corresponding PersistentVolumeClaim is deleted?

- A. Retain
- B. Recycle
- C. Delete
- D. Recreate

**Answer:** C

**Explanation:**

- **Delete (C)** removes the PersistentVolume when its PersistentVolumeClaim is deleted.
  - **Retain (A)** keeps the volume even after the claim is deleted.
  - **Recycle (B)** is deprecated in favor of external storage solutions.
  - **Recreate (D)** is not a valid reclaim policy.
-

## 5. Kubernetes etcd Backup and Restore

**Question:** What is the recommended method to back up an **etcd** database in a Kubernetes cluster?

- A. Use `kubectl get all -o yaml` and save the output.
- B. Take a snapshot using `etcdctl snapshot save`.
- C. Export the Kubernetes ConfigMap data.
- D. Run `kubectl backup cluster` (non-existent command).

**Answer:** B

**Explanation:**

- `etcdctl snapshot save` (B) creates a point-in-time backup of the etcd database.
  - **Kubectl YAML exports** (A) do not include all cluster states.
  - **ConfigMaps** (C) do not store full cluster state.
  - **Option D is an invalid command.**
- 

## 6. Kubernetes Pod Security Standards

**Question:** You need to **enforce least privilege security** by ensuring that pods do **not** run as root. How can you enforce this policy?

- A. Using NetworkPolicies
- B. Setting a SecurityContext in the Pod spec
- C. Using a Kubernetes ServiceAccount
- D. By enabling Pod Priority Classes

**Answer:** B

**Explanation:**

- **SecurityContext** (B) allows you to enforce pod security policies such as preventing root access.
  - **NetworkPolicies** (A) control network access but not security privileges.
  - **ServiceAccounts** (C) manage API authentication, not runtime security.
  - **Pod Priority Classes** (D) control scheduling priority, not security.
-

## 7. Kubernetes Audit Logs

**Question:** Where does Kubernetes store **API audit logs** by default?

- A. etcd database
- B. Systemd logs on worker nodes
- C. `/var/log/kube-apiserver-audit.log`
- D. `/var/lib/kubernetes/logs`

**Answer:** C

**Explanation:**

- Kubernetes audit logs are stored in `/var/log/kube-apiserver-audit.log` (C) when enabled.
  - etcd (A) stores cluster state, not logs.
  - Systemd logs (B) only log system-related events.
  - Path D is incorrect and does not exist by default.
- 

## 8. Kubernetes Node Drain

**Question:** What happens when you run `kubectl drain <node>`?

- A. All running pods on the node are forcibly deleted immediately.
- B. Pods with ReplicaSets are rescheduled on other nodes, while standalone pods are evicted.
- C. All DaemonSet pods are evicted.
- D. New pods can still be scheduled on the drained node.

**Answer:** B

**Explanation:**

- Drain (B) evicts pods managed by ReplicaSets but cannot evict DaemonSet pods unless explicitly specified.
  - Pods are not deleted immediately (A), they are gracefully evicted.
  - DaemonSet pods are not evicted by default (C).
  - New pods are prevented from scheduling (D is incorrect).
-

## 9. Kubernetes Pod Disruption Budget (PDB)

**Question:** You want to ensure that at least **one pod of a deployment** remains running during node maintenance. How can you enforce this?

- A. Set a minimum `replicas: 1` in the Deployment.
- B. Create a `PodDisruptionBudget (PDB)`.
- C. Use a `PriorityClass` to protect the pod.
- D. Use a Horizontal Pod Autoscaler.

**Answer:** B

**Explanation:**

- **PodDisruptionBudget (B)** ensures a minimum number of available replicas during voluntary disruptions.
  - **Replica count (A)** does not prevent disruptions.
  - **PriorityClasses (C)** affect scheduling priority, not disruptions.
  - **HPA (D)** scales based on metrics, not maintenance events.
- 

## 10. Kubernetes Secrets Security

**Question:** How can you **secure Kubernetes Secrets** stored in a cluster?

- A. Store secrets as ConfigMaps.
- B. Use `kubectl create secret` and enable encryption at rest.
- C. Pass secrets as environment variables in pod specs.
- D. Store secrets in plaintext inside manifests.

**Answer:** B

**Explanation:**

- **Option B is the correct approach—Kubernetes Secrets should be encrypted at rest using KMS or custom encryption providers.**
  - **ConfigMaps (A)** are not designed for sensitive data.
  - **Environment variables (C) expose secrets in pod specifications.**
  - **Plaintext in manifests (D) is insecure and violates security best practices.**
-

# SET D

## 1. Kubernetes Role-Based Access Control (RBAC)

**Question:** You need to allow a developer to **view and list pods** in a namespace but prevent them from deleting pods. Which RBAC resource should you create?

- A. **RoleBinding** with a **ClusterRole** of **admin**
- B. **ClusterRoleBinding** with the **view** role
- C. **Role** with **get**, **list**, and **watch** permissions on pods
- D. **RoleBinding** that grants **delete** permissions only for the developer's pod

**Answer:** C

**Explanation:**

- A **Role (C)** scoped to a namespace with **get**, **list**, and **watch** permissions is correct.
  - A **ClusterRoleBinding (B)** applies to the whole cluster, which is not needed here.
  - Option D grants delete permissions, which is what we want to avoid.
- 

## 2. Kubernetes Static Pods

**Question:** How do **static pods** differ from regular Kubernetes pods?

- A. **Static pods are created by the kube-apiserver, while regular pods are created by kubelet.**
- B. **Static pods are managed directly by the kubelet and do not have an associated Deployment or ReplicaSet.**
- C. **Static pods can be rescheduled to another node if the node fails.**
- D. **Static pods do not support logs or monitoring via `kubectl logs`.**

**Answer:** B

**Explanation:**

- **Static pods are managed by the kubelet (B) and do not go through the API server.**
  - **They are not rescheduled (C) because there is no controller managing them.**
  - **Logs can still be accessed using `kubectl logs` (D is incorrect).**
-

## 3. Kubernetes Ingress TLS Configuration

**Question:** You need to secure an **Ingress resource** with TLS encryption. What steps should you take?

- A. Enable HTTPS in the Ingress controller and provide a Secret containing the TLS certificate.
- B. Modify the Service definition to use TLS instead of HTTP.
- C. Store the TLS certificate in a ConfigMap and reference it in the Ingress spec.
- D. Configure NetworkPolicy to enforce TLS between pods.

**Answer:** A

**Explanation:**

- Option A is correct—you must enable HTTPS and provide a Kubernetes Secret containing the TLS certificate.
  - Ingress handles TLS, not the Service (B is incorrect).
  - ConfigMaps (C) are not used for TLS certificates.
  - NetworkPolicies (D) control traffic but do not enforce TLS at the ingress level.
- 

## 4. Kubernetes Node Selectors vs. Affinity

**Question:** What is the key difference between **nodeSelector** and **nodeAffinity**?

- A. nodeSelector supports expressions, while nodeAffinity only supports exact matches.
- B. nodeAffinity is more flexible and supports **preferred** and **required** scheduling rules, whereas nodeSelector only supports exact label matches.
- C. nodeSelector requires multiple labels, whereas nodeAffinity requires only one label.
- D. nodeAffinity does not allow node selection based on labels.

**Answer:** B

**Explanation:**

- nodeSelector (B) only supports exact label matches.
  - nodeAffinity (B) supports expressions and **preferred** vs. **required** rules.
  - nodeSelector does not require multiple labels (C is incorrect).
- 

## 5. Kubernetes Troubleshooting - CrashLoopBackOff

**Question:** A pod is stuck in `CrashLoopBackOff`. What is the best command to investigate the issue?

- A. `kubectl describe pod <pod-name>`
- B. `kubectl get events --all-namespaces`
- C. `kubectl logs <pod-name>`
- D. All of the above

**Answer:** D

**Explanation:**

- `kubectl describe pod` (A) provides detailed pod information and reasons for restarts.
  - `kubectl get events` (B) helps check for related issues in the namespace.
  - `kubectl logs` (C) retrieves logs for debugging the application itself.
- 

## 6. Kubernetes Network Plugins (CNI)

**Question:** What is the role of a **Container Network Interface (CNI)** in Kubernetes?

- A. To provide pod-to-pod networking across nodes
- B. To expose Kubernetes services externally
- C. To encrypt all network traffic between pods
- D. To configure TLS certificates for communication between nodes

**Answer:** A

**Explanation:**

- CNI plugins (A) handle pod networking within and across nodes.
  - Services (B) use different mechanisms, not CNI.
  - Encryption (C) can be handled by additional tools like Istio or service meshes.
  - TLS certificates (D) are not managed by CNI.
- 

## 7. Kubernetes Pod Eviction

**Question:** What happens when a node reaches its memory limit?

- A. The node shuts down automatically.
- B. Kubernetes evicts low-priority pods to free up memory.
- C. All pods on the node are restarted.
- D. Pods are automatically moved to another node.

**Answer:** B

**Explanation:**

- Kubernetes evicts pods (B) based on priority and QoS classes.
  - Nodes do not shut down automatically (A).
  - Pods are not restarted unless a controller handles them (C).
  - Pods are not "moved" but rather restarted elsewhere if scheduling allows (D).
- 

## 8. Kubernetes etcd Key-Value Store

**Question:** What type of data is stored in **etcd**?

- A. Pod logs
- B. Cluster configuration, API objects, and state data
- C. Docker images used by the nodes
- D. PersistentVolumeClaim data

**Answer:** B

**Explanation:**

- etcd stores the entire Kubernetes cluster state and configuration (B).
  - Pod logs (A) are stored separately in logging systems.
  - Docker images (C) are stored in container registries, not etcd.
  - PersistentVolumeClaims (D) reference storage, but the data itself is elsewhere.
- 

## 9. Kubernetes Helm Package Manager

**Question:** What is the purpose of **Helm** in Kubernetes?

- A. To manage application deployments using pre-defined templates
- B. To monitor Kubernetes logs and metrics
- C. To provide a network overlay for pod communication
- D. To perform automatic scaling of workloads

**Answer:** A

**Explanation:**

- Helm (A) helps package, install, and manage Kubernetes applications.
  - Logging (B) is handled by tools like Fluentd, Loki, and ELK.
  - Networking (C) is handled by CNI plugins.
  - Scaling (D) is handled by the HPA or Cluster Autoscaler.
- 

## 10. Kubernetes API Server Authentication

**Question:** Which method is NOT a valid authentication mechanism for the **Kubernetes API Server**?

- A. X.509 Client Certificates
- B. ServiceAccount Tokens
- C. NetworkPolicies
- D. OIDC (OpenID Connect) Tokens

**Answer:** C

**Explanation:**

- NetworkPolicies (C) control networking but do not handle authentication.
  - X.509 certificates (A), ServiceAccount tokens (B), and OIDC tokens (D) are all valid authentication methods.
- 

## SET E

### 1. Kubernetes Role-Based Access Control (RBAC)

**Question:** You need to grant a user permission to create, modify, and delete deployments in a specific namespace. Which resource should you use?

- A. ClusterRoleBinding
- B. RoleBinding with a Role that allows access to deployments
- C. ServiceAccount with elevated privileges
- D. NetworkPolicy that allows API access from the user's workstation

**Answer:** B

**Explanation:**

- **RoleBinding (B)** grants permissions within a specific namespace.
  - **ClusterRoleBinding (A)** applies cluster-wide, which is unnecessary here.
  - **ServiceAccounts (C)** are for workloads, not user access.
  - **NetworkPolicies (D)** do not control API access.
- 

## 2. Kubernetes Control Plane Components

**Question:** Which Kubernetes component is responsible for **scheduling pods onto nodes?**

- A. **kube-apiserver**
- B. **kubelet**
- C. **kube-controller-manager**
- D. **kube-scheduler**

**Answer:** D

**Explanation:**

- **kube-scheduler (D)** assigns pods to nodes based on resource availability and constraints.
  - **kube-apiserver (A)** is the front-end of the control plane.
  - **kubelet (B)** runs on each node and manages pods.
  - **kube-controller-manager (C)** manages controllers like ReplicaSet and Node Controller.
- 

## 3. Kubernetes Network Policies

**Question:** You need to block **all ingress traffic** to a set of pods except from a specific namespace. What should you use?

- A. **Ingress Controller**
- B. **NetworkPolicy with `Ingress` rules**
- C. **NodeSelector to restrict pod placement**
- D. **ServiceAccount-based authentication**

**Answer:** B

**Explanation:**

- NetworkPolicy (B) controls pod-to-pod traffic based on namespace and label selectors.
  - Ingress Controllers (A) handle external traffic.
  - NodeSelector (C) controls pod scheduling but does not block traffic.
  - ServiceAccounts (D) manage API access, not network traffic.
- 

## 4. Kubernetes High Availability

**Question:** How do you make an **etcd** cluster highly available in a Kubernetes control plane?

- A. Run a single etcd instance on the master node.
- B. Deploy etcd in a multi-node configuration with an odd number of instances.
- C. Use a PersistentVolume to store etcd data on a single node.
- D. Configure a network policy that ensures only kube-apiserver can communicate with etcd.

**Answer:** B

**Explanation:**

- etcd should be deployed with an odd number of nodes (B) to maintain quorum.
  - Single-node etcd (A) is a single point of failure.
  - PersistentVolumes (C) do not provide HA.
  - NetworkPolicies (D) enhance security but do not affect HA.
- 

## 5. Kubernetes Readiness Probe

**Question:** What happens when a **readiness probe** fails for a running pod?

- A. The pod is restarted immediately.
- B. The pod is removed from the Service's endpoints but continues running.
- C. The pod is evicted from the node.
- D. All containers in the pod are automatically stopped.

**Answer:** B

**Explanation:**

- A failing readiness probe (B) removes the pod from service load balancers but does not restart it.
- Liveness probes (A) trigger restarts, not readiness probes.

- Pods are not evicted (C) when readiness fails.
  - Containers remain running (D).
- 

## 6. Kubernetes Persistent Volume Claims

**Question:** What happens when a **PersistentVolumeClaim (PVC)** is deleted but the underlying **PersistentVolume (PV)** uses the **Retain** reclaim policy?

- A. The **PersistentVolume** is deleted automatically.
- B. The **PersistentVolume** is detached but remains available for manual recovery.
- C. The **PersistentVolume** is recycled for reuse by other PVCs.
- D. All data on the **PersistentVolume** is erased immediately.

**Answer:** B

**Explanation:**

- Retain (B) keeps the **PersistentVolume** after PVC deletion for manual reuse.
  - Delete (A) is a different reclaim policy that removes the PV.
  - Recycle (C) is deprecated.
  - Data is not erased automatically (D).
- 

## 7. Kubernetes Pod Priority and Preemption

**Question:** A high-priority pod cannot start due to insufficient resources, but there are lower-priority pods running. What will Kubernetes do?

- A. Wait indefinitely until more nodes are added.
- B. Evict lower-priority pods to free up resources.
- C. Restart the lower-priority pods with reduced resource limits.
- D. Ignore priority settings and schedule the pod when possible.

**Answer:** B

**Explanation:**

- Kubernetes preemptively evicts lower-priority pods (B) to make room.
- It does not wait indefinitely (A).
- Pods are not restarted with modified limits (C).
- Priority is respected (D is incorrect).

---

## 8. Kubernetes API Aggregation Layer

**Question:** What is the purpose of the **API Aggregation Layer** in Kubernetes?

- A. To expose additional APIs without modifying the core Kubernetes API server.
- B. To increase the performance of API requests by caching responses.
- C. To enforce security policies for external API calls.
- D. To provide an alternative to etcd for storing API objects.

**Answer:** A

**Explanation:**

- API Aggregation (A) allows new APIs to be added without modifying the API server.
  - It does not cache responses (B).
  - Security policies (C) are handled elsewhere.
  - etcd remains the primary data store (D).
- 

## 9. Kubernetes CRI (Container Runtime Interface)

**Question:** Which container runtime is NOT compatible with Kubernetes' **Container Runtime Interface (CRI)**?

- A. containerd
- B. Docker (without CRI plugin)
- C. CRI-O
- D. gVisor

**Answer:** B

**Explanation:**

- Docker (B) is not CRI-compatible by default and requires the `cri-dockerd` plugin.
  - containerd (A) and CRI-O (C) are CRI-compliant.
  - gVisor (D) is a sandbox but integrates with CRI runtimes.
- 

## 10. Kubernetes Troubleshooting - API Server

**Question:** The **kube-apiserver** is down. What command can you run on a control plane node to check its status?

- A. `kubectl get pods -n kube-system`
- B. `systemctl status kube-apiserver`
- C. `kubectl logs kube-apiserver`
- D. `etcdctl cluster-health`

**Answer:** B

**Explanation:**

- `systemctl status kube-apiserver` (B) checks the API server service status.
- `kubectl` (A, C) depends on the API server, which is down.
- `etcdctl` (D) checks etcd, not the API server.